

A Survey of Stream Processing Languages

Martin Hirzel
IBM Research AI, USA
hirzel@us.ibm.com

Emanuele Della Valle
Politecnico di Milano, Italy
emanuele.dellavalle@polimi.it

Guillaume Baudart
IBM Research AI, USA
Guillaume.Baudart@ibm.com

Sherif Sakr
KSAU-HS and UNSW
ssakratcse.unsw.edu.au

Angela Bonifati
Lyon 1 University, France
angela.bonifati@univ-lyon1.fr

Akrivi Vlachou
University of Piraeus, Greece
avlachou@aueb.gr

```
1 SELECT IStream(Max(len) AS mxl,  
2           MaxCount(len) AS num,  
3           ArgMax(len, caller) as who)  
4 FROM Calls[Range 24 Hours Slide 1 Minute]
```

Figure 1: CQL code example.

- CEP (complex event processing) / MATCH-RECOGNIZE [18] [11]
- XML / NiagaraCQ [9], YFilter [10] [Sherif](#)
- RDF / C-SPARQL [6] [Emanuele](#), [Akrivi](#)
- stream reasoning [Emanuele](#)
- reactive spreadsheets / ActiveSheets [17] [Martin](#)
- controlled natural language / META [5] [Martin](#)
- close with summary/comparison table
- perhaps according to performance/generalizability/productivity

ABSTRACT

TODO

1. INTRODUCTION

TODO ~0.8 pages

- motivation
 - importance and rise of streaming
 - benefits of languages
 - diversity of languages, lack of standard
 - lack of a recent survey [16] [14]
- background / definitions
 - stream, streaming application, DSL, stream processing language
 - EDSLs [13], recent trends in EDSLs
 - requirements: performance, generality, productivity
 - most streaming languages are declarative, so the traditional paradigm categories imperative or functional don't apply
- roadmap for rest of paper
- take-home message: democratization of streaming

2. STREAM PROCESSING LANGUAGES

TODO ~2.8 pages total

- this section will have ~0.3 pages about each language, and each of these snippets will be structured around questions: why-who-when-what-where-when (see CQL for an example, but the others will be similar!)
- before the snippets on each of the languages, we will briefly introduce and explain these questions, as follows:
 - why: objective, audience, domain
 - who: inventors, supporters
 - when: first release / first paper
 - what: key idea, data model, type system, code example
 - where: being developed or offered today, license
 - whence: influenced-by and influences
- descriptions of individual languages
 - relational / CQL [2] [Sherif](#)
 - why: precise semantics [4], example app linear-road [3]
 - who: Arasu/Widom at Stanford
 - when: 2004
 - what: algebra of R2R (relational), S2R (windows), R2S (I/R/D); data model relational; type system [15]; example Figure 1
 - where: Stanford STREAM no longer active
 - whence: influenced-by TelegraphCQ [8], influences StreamInsight [1] and many others
 - synchronous dataflow / Lustre [7] [Guillaume](#)
 - big-data streaming / SPL [12] [Martin](#), [Emanuele](#)

3. WHAT'S NEXT?

TODO ~0.6 pages total

- this section will have ~0.2 pages about each challenge, where the challenges come from our Dagstuhl discussion, and each snippet will be structured as follows:
 - what's the problem
 - why is solving it useful
 - why is it hard
 - what makes it a streaming languages problem
- descriptions of individual challenges:
 - Handling data variety while keeping the language simple and fast [Emanuele](#)
 - Handling veracity in a simple and well-defined way [Akrivi](#)
 - Getting broad adoption for a streaming language

4. CONCLUSION

TODO ~0.2 pages

Acknowledgements

The idea for this paper was conceived at Dagstuhl Seminar 17441 on “Big Stream Processing Systems”.

5. REFERENCES

- [1] M. H. Ali, C. Gere, B. Raman, B. Sezgin, T. Tarnavski, T. Verona, P. Wang, P. Zabback, A. Kirilov, A. Ananthanarayan, M. Lu, A. Raizman, R. Krishnan, R. Schindlauer, T. Grabs, S. Bjeletich, B. Chandramouli, J. Goldstein, S. Bhat, Y. Li, V. Di Nicola, X. Wang, D. Maier, I. Santos, O. Nano, and S. Grell. Microsoft CEP server and online behavioral targeting. In *Demo at Very Large Data Bases (VLDB-Demo)*, pages 1558–1561, 2009.
- [2] A. Arasu, S. Babu, and J. Widom. The CQL continuous query language: semantic foundations and query execution. *Journal on Very Large Data Bases (VLDB J.)*, 15(2):121–142, 2006.
- [3] A. Arasu, M. Cherniack, E. Galvez, D. Maier, A. S. Maskey, E. Ryvkina, M. Stonebraker, and R. Tibbetts. Linear road: A stream data management benchmark. In *Conference on Very Large Data Bases (VLDB)*, pages 480–491, 2004.

- [4] A. Arasu and J. Widom. A denotational semantics for continuous queries over streams and relations. *SIGMOD Record*, 33(3), Sept. 2004.
- [5] M. Arnold, D. Grove, B. Herta, M. Hind, M. Hirzel, A. Iyengar, L. Mandel, V. Saraswat, A. Shinnar, J. Siméon, M. Takeuchi, O. Tardieu, and W. Zhang. META: Middleware for events, transactions, and analytics. *IBM Journal of Research and Development*, 60(2-3):15:1–15:10, 2016.
- [6] D. F. Barbieri, D. Braga, S. Ceri, E. Della Valle, and M. Grossniklaus. C-SPARQL: SPARQL for continuous querying. In *Poster at International World Wide Web Conferences (WWW-Poster)*, pages 1061–1062, 2009.
- [7] P. Caspi, D. Pilaud, N. Halbwachs, and P. Raymond. LUSTRE: a declarative language for real-time programming. In *Symposium on Principles of Programming Languages (POPL)*, pages 178–188, 1987.
- [8] S. Chandrasekaran, O. Cooper, A. Deshpande, M. J. Franklin, J. M. Hellerstein, W. Hong, S. Krishnamurthy, S. Madden, V. Raman, F. Reiss, and M. A. Shah. TelegraphCQ: continuous dataflow processing for an uncertain world. In *Conference on Innovative Data Systems Research (CIDR)*, 2003.
- [9] J. Chen, D. J. DeWitt, F. Tian, and Y. Wang. NiagaraCQ: A scalable continuous query system for internet databases. In *International Conference on Management of Data (SIGMOD)*, pages 379–390, 2000.
- [10] Y. Diao, P. M. Fischer, M. J. Franklin, and R. To. YFilter: Efficient and scalable filtering of XML documents. In *Demo at International Conference on Data Engineering (ICDE-Demo)*, pages 341–342, 2002.
- [11] M. Hirzel. Partition and compose: Parallel complex event processing. In *Conference on Distributed Event-Based Systems (DEBS)*, pages 191–200, 2012.
- [12] M. Hirzel, S. Schneider, and B. Gedik. SPL: An extensible language for distributed stream processing. *Transactions on Programming Languages and Systems (TOPLAS)*, 39(1):5:1–5:39, March 2017.
- [13] P. Hudak. Modular domain specific languages and tools. In *International Conference on Software Reuse (ICSR)*, pages 134–142, 1998.
- [14] W. M. Johnston, J. R. P. Hanna, and R. J. Millar. Advances in dataflow programming languages. *ACM Computing Surveys (CSUR)*, 36(1):1–34, 2004.
- [15] R. Soulé, M. Hirzel, B. Gedik, and R. Grimm. River: An intermediate language for stream processing. *Software – Practice and Experience (SP&E)*, 46(7):891–929, July 2016.
- [16] R. Stephens. A survey of stream processing. *Acta Informatica*, 34(7):491–541, 1997.
- [17] M. Vaziri, O. Tardieu, R. Rabbah, P. Suter, and M. Hirzel. Stream processing with a spreadsheet. In *European Conference on Object-Oriented Programming (ECOOP)*, pages 360–384, 2014.
- [18] F. Zemke, A. Witkowski, M. Cherniak, and L. Colby. Pattern matching in sequences of rows. Technical report, ANSI Standard Proposal, 2007.